

Are all the $\Delta\theta$'s computed by the nodes every ~15 seconds stored locally in those nodes to be made available for the epoch calculation every 30 days? "the current θ hash is effectively committed to the on-chain consensus state" - how does this actually work from an implementation perspective? Does commitment mean it is stored in the mempool or DHT? Where precisely is it committed to or stored?

1. Are all the $\Delta\theta$'s computed by the nodes every ~15 seconds stored locally in those nodes to be made available for the epoch calculation every 30 days?

No – the $\Delta\theta$ (gradient updates) are not stored locally on nodes for the full 30-day epoch. Instead, they are:

- Computed locally inside the node's TEE every ~15 seconds or every 1000 transactions.
- Encrypted and immediately submitted to a secure aggregation pool inside a TEE cluster.
- Aggregated intra-shard using secure multi-party computation (SMPC) within seconds.
- The aggregated gradient is then used to update the global encoder θ incrementally via federated learning.
- Only the final updated θ (after 30 days) is committed to the chain; individual $\Delta\theta$'s are discarded after secure aggregation to preserve privacy and minimize storage.

This ensures no single node retains a history of gradients, aligning with NERV's privacy and security model (TEE-bound, DP-SGD).

2. How is the current θ hash committed to the on-chain consensus state from an implementation perspective?

The θ hash (hash of the neural encoder weights) is committed as part of the embedding root in the consensus state. Specifically:

- After each batch execution, the new embedding root $e_{\{t+1\}}$ is computed and hashed to $h_{\{t+1\}}$.
- This root includes the current θ hash, because the embedding $e_{\{t+1\}}$ is a function of the state S_t and the encoder θ .
- During AI-native consensus, validators agree on the new embedding root (which implicitly attests to the correct θ).
- The root is then broadcast via DHT and gossip and stored in the canonical shard lattice – a DAG of embedding roots that represents the global state.

- Light clients sync only the chain of embedding root hashes (32 bytes each), which serves as the commitment to both state and encoder.

In short: the θ hash is not stored separately in a mempool or DHT alone; it is embedded within the cryptographically committed embedding root, which is validated by consensus and recorded in the shard lattice. This ensures verifiable and tamper-proof commitments to the encoder version used in each state transition.

Where does the canonical shard lattice — a DAG of embedding roots reside? What is a DAG? Also, where is the aggregated gradient θ stored (because the nodes provide $\Delta\theta$'s every ~15 seconds and the epoch only uses the aggregated gradients every 30 days for the new public θ' to be computed, it needs to be stored somewhere for the epoch to access the aggregated gradients over the last 30 days)?

1. What is a DAG, and where does the canonical shard lattice (DAG of embedding roots) reside?

What is a DAG?

DAG stands for Directed Acyclic Graph.

In simple terms:

- It's a graph where edges have a direction (like parent $\rightarrow$ child), and no cycles exist (you can't follow edges and return to the same node).
- In NERV, each embedding root is a node in the DAG, and edges represent transitions from one state root to the next (per shard and cross-shard dependencies).

Where does the shard lattice reside?

The canonical shard lattice is stored across all full nodes in a decentralized, peer-to-peer manner, not in a central database. Specifically:

- In each full node's memory/SSD: Every full node maintains a copy of the DAG structure locally.

- Gossiped and synchronized via libp2p/IPFS-like protocols: When a new embedding root is finalized, it's broadcast via gossip; nodes update their local DAG.
- Light clients only store the chain of hashes (the “spine” of the DAG), not the full DAG.
- The DAG itself is content-addressed — each embedding root is referenced by its hash (BLAKE3), enabling distributed storage (like IPFS).

In essence: The lattice is a distributed, immutable data structure replicated across the network, not stored in a single location.

2. Where is the aggregated gradient θ stored between the 15-second updates and the 30-day epoch?

This is a subtle but important point. Let's clarify the workflow:

Step-by-step gradient flow in NERV:

1. Every ~15 seconds, each node computes a local gradient $\Delta\theta_i$ inside its TEE.
2. Immediate secure aggregation (in TEE clusters) happens intra-shard:
 - Gradients are summed using SMPC.
 - The result is an updated global model θ_{new} , not stored per-node gradients.
 - The aggregation is incremental — think of it as online learning.
3. This updated θ_{new} is used immediately for the next batch of transactions.
4. No long-term storage of raw or aggregated gradients is needed — the model updates continuously.

So where is the “aggregated gradient” stored for the 30-day epoch?

The epoch update does not re-use past gradients. Instead:

- Every 30 days, a federated learning round is run to produce a new encoder θ' .

- This round uses fresh, recent transaction data (from the last epoch) to retrain/tune the encoder.
- The training happens inside TEEs with secure aggregation, just like the 15-second updates, but with a larger dataset and more rounds.
- The resulting θ' is then committed to the chain (via the embedding root), and nodes switch to it.

Why no long-term gradient storage?

- Privacy: Storing gradients long-term could leak information.
- Efficiency: Federated learning in NERV is continuous and incremental; the model is always up-to-date.
- Epoch training starts from the current θ and fine-tunes it using recent data — no need to replay old gradients.

In summary:

The aggregated gradient isn't stored for 30 days. The model θ evolves continuously via TEE-secured federated learning. The 30-day epoch is a scheduled retraining checkpoint that produces a new encoder θ' using recent data, not stored gradients.

Summary Table

Concept	What it is	Where it resides
Shard Lattice (DAG)	Directed graph of embedding roots (state hashes)	Distributed across all full nodes, gossiped and stored locally

Aggregated Gradients ($\Delta\theta$)	Incremental model updates from FL	Not stored long-term; model updates continuously in TEEs
---	-----------------------------------	--

Encoder θ	Current neural encoder weights	Committed in embedding roots; full weights distributed via DHT
------------------	--------------------------------	--

Epoch θ'	New encoder every 30 days	Produced via fresh FL, committed to chain, distributed to nodes
-----------------	---------------------------	---

Please map out the data flow diagram for gradient aggregation and epoch transitions. Also, the nodes use anonymized recent transaction batches for computing the local $\Delta\theta$'s every ~15 seconds. Does the epoch also train on recent anonymized recent transaction batches for computing the new θ' ? If so, what's the difference in the number of the recent transactions between the nodes' useful work and the epoch's θ' calculation?


Data Flow Diagram: Gradient Aggregation & Epoch Transitions


text

NERV GRADIENT FLOW & EPOCH TRANSITION ARCHITECTURE


=====

[CONTINUOUS USEFUL-WORK CYCLE (Every ~15 seconds)]





[EPOCH TRANSITION (Every 30 days)]



Dataset size: ~10-100× larger than 15-second batches	
Phase 2: Multi-Round Federated Learning	
For round = 1 to R (R ~ 50-200):	
1. Nodes train on local epoch dataset → $\Delta\theta^{(r)}_i$	
2. Secure aggregation in TEE cluster → $\Delta\theta^{(r)}_{\text{shard}}$	
3. $\theta \leftarrow \theta + \eta \cdot \Delta\theta^{(r)}_{\text{shard}}$ (update global model)	
4. Check homomorphism preservation (error < 1e-9)	
Phase 3: Finalization & Commitment	
1. Final θ' validated via Halo2 proof:	
• Homomorphism preserved	
• Performance improved	
2. θ' hash committed in special "epoch transition block"	
3. Full θ' weights distributed via erasure-coded DHT	
4. All nodes switch to θ' at next block	

[STORAGE LOCATIONS & DATA RETENTION]

Data Type	Storage Location	Retention Period
Individual $\Delta\theta_i$ (15-second)	Never stored (TEE volatile memory)	Destroyed after aggregation (~15 seconds)
Aggregated $\Delta\theta$ (shard-level)	TEE-sealed memory (ephemeral)	Until next aggregation (~seconds)
Current θ (global model)	- TEE memory (active) - DHT (redundant) - Embedding root hash	Until epoch transition (30 days max)
Epoch Dataset (for θ' train)	TEE-sealed encrypted cache (LRU eviction)	14-30 days (rolling window)
Historical θ' (past encoders)	- Arweave/IPFS - Light client headers	Permanent (for verification)

Transaction Volume Comparison: Useful-Work vs Epoch Training

Useful-Work (15-second cycles):

- Volume per node: ~100-1,000 transactions per cycle
- Source: Most recent anonymized batches (last few minutes)
- Purpose: Incremental model refinement, adaptability to current trends
- Training: Single gradient step (SGD), emphasis on recent patterns

Epoch Training (30-day cycles):

- Volume per node: ~500,000-5,000,000+ transactions
- Source: Stratified sample from entire epoch (30 days)
- Purpose: Comprehensive model retraining, learning long-term patterns
- Training: 50-200 full training rounds, emphasis on stability/generalization

Numerical Example:

text

Assume: 1M TPS network × 30 days × 86,400 seconds/day

= 2.592 trillion transactions per epoch

Per node sampling (0.0001%): ~2.6 million transactions for epoch training

vs 15-second cycle: ~1,500 transactions (at 100 TPS per shard)

Ratio: Epoch dataset ~1,700× larger than 15-second batch

Key Differences Between the Two Training Processes

Aspect	Useful-Work (15s)	Epoch Training (30d)
Frequency	Continuous (~15s)	Periodic (30 days)
Data Recency	Last few minutes	Full epoch (30 days)
Training Intensity	Single gradient step	50-200 full rounds
Purpose	Adapt to immediate trends	Learn stable patterns
Storage	No gradient retention	Cached dataset for epoch
Shapley Rewards	Yes (per cycle)	No (built into epoch update)
Model Impact	Incremental θ adjustment	Complete θ' replacement

Critical Implementation Details

Why Epoch Training Needs Different Data:

1. Statistical Significance: Need enough data to ensure homomorphism preservation proof holds
2. Diversity: Must capture various transaction patterns over time
3. Stability: Prevents overfitting to recent anomalies

Privacy Preservation in Epoch Training:

- Epoch datasets are not stored as raw transactions
- Instead, nodes maintain TEE-sealed feature vectors from recent batches
- Training happens on anonymized embeddings, not plaintext data
- Differential privacy ($\sigma=0.5$) applied in both phases

Epoch Transition Mechanics:

1. At epoch boundary, a governance trigger initiates retraining
2. Nodes contribute to secure multi-round federated learning
3. Final θ' validated via Halo2 proof of homomorphism preservation
4. New encoder activated via hard fork protocol (automated, non-contentious)

Summary: The system maintains a continuous learning loop (15-second cycles) for adaptability, plus periodic deep retraining (30-day epochs) for stability. The epoch training uses $\sim 1,000\times$ more data than individual useful-work cycles, enabling comprehensive model refinement while preserving privacy through TEEs and differential privacy.

How is the epoch dataset arrived at (the subset of sample from a huge sample space)? Also, the dataset contains TEE-sealed feature vectors - where and how are those vectors stored and sampled for the epoch calculation?

Epoch Dataset Construction & Feature Vector Storage


How the Epoch Dataset is Constructed

text


EPOCH DATASET CONSTRUCTION PIPELINE

=====

[RAW TRANSACTION STREAM]



[PER-NODE PROCESSING (TEE)]



| | Output: $FV(B) \in \mathbb{R}^{1024}$ | |

| | (1024-dimensional anonymized vector) | |

| | | |

| | | |

| 3. Cache Management: |

| • $FV(B)$ encrypted with ML-KEM |

| • Stored in TEE-sealed LRU cache |

| • Max: 100,000 vectors per node (~100MB) |

| • Eviction: Oldest when full |

| | |

| |

[EPOCH TRIGGER (Day 30)]

| |

| ▼

| |

| Stratified Sampling for Epoch Dataset: |

| |

| 1. Temporal Stratification: |

| • Divide 30 days into 720 time windows (1-hour) |

| • Sample equal # vectors from each window |

| |

| 2. Shard Stratification: |

| • Ensure representation from all active shards |

| |

3. Volume-Based Weighting:	
• Batches with more txs get higher probability	
• But capped to prevent bias	
4. Final Dataset per Node:	
• Size: 50,000-100,000 feature vectors	
• Represents ~50M-100M original transactions	
(0.02-0.04% of total epoch txs)	

Where and How Feature Vectors Are Stored

Storage Architecture:

text

TEE-SEALED FEATURE VECTOR STORAGE

Node's TEE Memory Hierarchy:	
1. L0: Hot Cache (RAM, 10,000 vectors)	
• Most recent 2-3 days	
• Directly accessible for 15-second training	
2. L1: Encrypted NVMe Cache (100,000 vectors)	
• Full 30-day rolling window	

• Encrypted with TEE-bound keys
• LRU eviction when full
3. L2: Erasure-Coded Archive (Optional)
• Historical vectors (beyond 30 days)
• Reed-Solomon (k=10, m=4) distributed storage
• For research/audit only, not used in training

Feature Vector Format:

```
rust
struct FeatureVector {
    // Metadata (encrypted)

    timestamp: u64,           // When batch was processed
    shard_id: u32,            // Which shard
    batch_size: u16,          // Number of txs in batch
    dp_epsilon: f32,          // Privacy budget used

    // Data (encrypted)

    embeddings: [f32; 1024], // 1024D anonymized features
    delta_norm: f32,         // Magnitude of changes
    attention_weights: [f32; 256], // Aggregated attention patterns
}

// Total size: ~8.5KB per vector
```



```
// Encrypted with ML-KEM-768 + AES-256-GCM
```

Sampling Algorithm for Epoch Training

Stratified Reservoir Sampling:

```
python
```

```
def sample_epoch_dataset(tee_cache, target_size=50000):  
    """  
    Selects representative feature vectors for epoch training  
    while preserving privacy and temporal distribution  
    """  
  
    # Step 1: Temporal stratification  
    hourly_buckets = partition_by_hour(tee_cache, days=30)  
  
    samples = []  
    vectors_per_hour = target_size // 720 # ~69 vectors/hour  
  
    for hour_bucket in hourly_buckets:  
        if len(hour_bucket) == 0:  
            continue  
  
    # Step 2: Weight by batch diversity  
    weights = compute_diversity_score(hour_bucket)
```

```
# Step 3: DP sampling ( $\epsilon=0.1$  additional privacy)
```

```
selected = exponential_mechanism(
```

```
    candidates=hour_bucket,
```

```
    weights=weights,
```

```
    k=vectors_per_hour,
```

```
    epsilon=0.1
```

```
)
```

```
samples.extend(selected)
```

```
# Step 4: Ensure minimum from each shard
```

```
enforce_shard_minimum(samples, min_per_shard=100)
```

```
return samples[:target_size] # Final epoch dataset
```

Diversity Scoring Function:

```
python
```

```
def compute_diversity_score(feature_vector):
```

```
    """
```

```
    Scores how "interesting" a batch is for training
```

```
    Higher score = more valuable for encoder improvement
```

```
    """
```

```
    score = 1.0
```

```
# 1. Batch size factor (log scale)
```

```
score *= math.log(1 + feature_vector.batch_size)
```

```
# 2. Delta magnitude (measure of state change)
```

```
score *= 1 + feature_vector.delta_norm
```

```
# 3. Attention entropy (pattern diversity)
```

```
entropy = shannon_entropy(feature_vector.attention_weights)
```

```
score *= 1 + entropy
```

```
# 4. Temporal recency (recent batches slightly favored)
```

```
age_hours = (now - feature_vector.timestamp) / 3600
```

```
recency_factor = math.exp(-age_hours / 168) # 1-week half-life
```

```
score *= 1 + 0.1 * recency_factor
```

```
return score
```

Epoch Training Dataset vs Useful-Work Dataset

Characteristic	Useful-Work Dataset (15s)	Epoch Training Dataset
Source	Last 15 seconds of txs	Stratified 30-day window

Size	~1,000 txs	~50M txs (represented)
Storage	Volatile TEE memory	Encrypted TEE cache
Selection	Sequential (FIFO)	Stratified reservoir sampling
Privacy	$\epsilon=1.0$ DP	Additional $\epsilon=0.1$ for sampling
Purpose	Immediate adaptation	Long-term generalization

Critical Privacy Guarantees

1. Feature Vectors Don't Leak Individual Transactions:
 - o Each vector aggregates 100-1000 transactions
 - o DP noise added during extraction ($\epsilon=1.0$)
 - o Cannot reconstruct specific tx from vector
2. Sampling Preserves Distribution, Not Identities:
 - o Stratified sampling ensures temporal coverage
 - o Exponential mechanism adds DP to selection
 - o No single batch disproportionately influences
3. TEE-Bound Encryption:
4. rust

// Feature vectors are always encrypted

```
let sealed_vec = tee_seal!(
```

```
    key: per_node_attested_key,
```

```
data: feature_vector,  
  
policy: "LRU_30d"
```

5.);

Storage Requirements & Optimization

Component	Size per Node	Encryption	Retention
Hot cache (L0)	~85MB	TEE memory	2-3 days
Main cache (L1)	~850MB	ML-KEM + AES	30 days
Historical (L2)	Optional	Erasured-codes	Permanent

Optimizations:

- Quantization: 16-bit floats → 50% space reduction
- Deduplication: Similar vectors merged (cosine similarity >0.95)

- Compression: Zstd compression (2:1 ratio)
- Selective Archiving: Only 1% of vectors kept beyond 30 days

Epoch Training Workflow with Stored Vectors

text

EPOCH TRAINING EXECUTION

Phase 1: Dataset Preparation (Day 29-30)
• Each node runs sampling algorithm
• Produces local dataset D_local
• Verifies DP guarantees
Phase 2: Federated Learning Rounds (50-200 rounds)
For each round:
1. Node trains on D_local → gradient $\Delta\theta^{(r)}_i$
2. Secure aggregation in TEE cluster
3. Global model update
4. Homomorphism check (error < 1e-9)
Phase 3: Final Validation & Commitment
• Halo2 proof of homomorphism preservation
• θ' hash committed to chain
• Old cache cleared (except 1% for audit)

Key Insight: The epoch dataset isn't a raw transaction dump—it's a carefully constructed, privacy-preserving sample of feature vectors that capture transaction patterns without revealing specifics. This enables effective encoder retraining while maintaining NERV's privacy guarantees.

1. How is `train_step(θ, batch)` accomplished exactly?

This is the core of NERV's *Useful-Work* economy, detailed in Sections 4.3 and the Shapley-Value explanation (pages 26-29).

Process:

1. Input: A node receives a recent *anonymized* batch of transactions (`batch`) via the 5-hop TEE mixer. The node never sees plaintext amounts, sender/receiver IDs, or balances. It only sees the structured, encrypted data needed for gradient computation.
2. Objective: The node's task is to compute a gradient update ($\Delta\theta_i$) for the global neural encoder θ (the 24-layer transformer E_θ). The goal is to improve the encoder's ability to produce embeddings where the Transfer Homomorphism property holds with minimal error ($\leq 1e-9$).
3. Computation Inside TEE:
 - The node loads the current public encoder weights θ into its Trusted Execution Environment (TEE).
 - Using the anonymized batch as input, it performs a forward pass through E_θ to generate the predicted embedding e_{pred} .
 - It then computes a loss function. This loss likely measures the deviation from the perfect homomorphism. For example:
$$\text{Loss} = || (E_\theta(S_t) + \sum \delta(tx)) - E_\theta(S_{\{t+1\}}) ||^2$$
Where $S_{\{t+1\}}$ is the *implied* next state after applying the batch, inferred from the transaction deltas themselves.
 - A backward pass (backpropagation) is performed inside the TEE to calculate the gradient of this loss with respect to the encoder weights θ . This results in $\Delta\theta_i$.
 - Differential Privacy (DP) Noise: Before leaving the TEE, Gaussian noise ($\sigma=0.5$) is added to $\Delta\theta_i$. This ensures that the gradient update does not leak information about the specific private transactions in the batch (Page 27, DP-SGD reference).

4. Output: The noisy, encrypted gradient $\Delta\theta_i$ is sent to the intra-shard secure aggregator. The raw batch data is discarded from TEE memory.

In essence: `train_step( $\theta$ , batch)` is a standard ML training step (forward pass, loss calculation, backward pass) performed inside a privacy-hardened TEE on cryptographically blinded data, with DP noise applied to the output gradient.

2. Is the Global Model θ shard-specific? Is it stored implicitly?

This is a crucial point of understanding NERV's architecture.

- Global vs. Shard-Specific: The whitepaper indicates there is a global encoder θ (Page 33, "The encoder E_θ is fixed for an epoch...fully public and auditable"). This same encoder is used across all shards to ensure a consistent, homomorphic embedding space. The *unique state* of each shard is not stored in different θ values, but in their unique 512-byte embedding vector e_t .
 - What is Stored in the DAG Lattice: You are correct. The lattice (the canonical chain of state) stores the embedding root hashes ($h_t = \text{Hash}(e_t)$), not the weights θ . The weights θ are public parameters, updated every 30-day epoch and distributed via DHT/IPFS. The *state* is the compact vector e_t living in that shared, global embedding space defined by θ .
 - Implicit Calculation: The updated e_t (and thus the new root) is the direct result of applying the homomorphic delta in the execution phase: $e_{\{t+1\}} = e_t + \Delta_B$. The role of the updated global model θ' (from federated learning) is to be used in the *next epoch* to ensure future embeddings ($e_{\{t+500\}}$, etc.) continue to obey the linear homomorphism property with high fidelity.
-

3. How does the client-side wallet compute $\delta(tx)$?

This is the function of the Δ_θ sub-circuit (Page 21-22, "Transfer Homomorphism: A Deeper Explanation").

Process:

1. Fixed Public Tools: The user's wallet possesses two public, immutable pieces of code/data:
 - The current epoch's encoder circuit E_θ (or a lightweight "delta generator" Δ_θ derived from it).
 - The public account/identity embeddings. In the neural space, each account is mapped to a unique, fixed 512-dim vector (like a "word embedding" in an LLM). Sender (emb_s) and receiver (emb_r) are known to the wallet (as private secrets) but have public, verifiable embeddings in the circuit.
2. Local Computation:
 - To create a private transfer of amount amt from Alice to Bob, the wallet runs the Δ_θ sub-circuit *locally*.
 - The circuit takes as private inputs: emb_s, emb_r, amt .
 - It outputs the delta vector $\delta(tx)$. The whitepaper suggests (Page 21) this is computed approximately as:

$$\delta(tx) \approx amt * (emb_r - emb_s) + bias_terms$$
 - This $\delta(tx)$ is a 512-element vector that, when added to the current global embedding e_t , will correctly represent the state where amt has moved from emb_s to emb_r *without ever revealing* $amt, emb_s, or emb_r$.
3. Proof Generation: The wallet simultaneously generates a Halo2 ZK proof attesting that:
 - The computed $\delta(tx)$ is valid (follows the rules of the Δ_θ circuit).
 - The sender's hidden balance is sufficient (this is proven against a private balance commitment, not revealed).

This proof is wrapped inside the onion for submission.

In short: The wallet uses the public, epoch-frozen neural circuit (Δ_θ) and the known neural embeddings of the parties to *algebraically compute* the precise change vector $\delta(tx)$ in the high-dimensional latent space. This is the "magic" of the pre-trained homomorphism.

4. What is the difference between $e_{\{t+1\}} = e_t + \Delta_B$ and $e_{\{t+1\}} = e_t + \text{aggregated } \Delta\theta_shard$?

These are fundamentally different operations and should not be conflated.

- $e_{\{t+1\}} = e_t + \Delta_B$ (State Update - Page 9):
 - Δ_B is the sum of transaction deltas: $\sum \delta(tx_i)$ for all tx_i in batch B.
 - This is how the ledger state is updated. It's an operation in the embedding space (R^{512}), moving the state vector e_t to its new position $e_{\{t+1\}}$.
- $\theta' = \theta + \text{aggregated } \Delta\theta_{\text{shard}}$ (Model Update - Pages 26-27):
 - aggregated $\Delta\theta_{\text{shard}}$ is the average of gradient updates from nodes, used in Federated Learning.
 - This is how the encoder model parameters (weights θ) are updated. It's an operation in the parameter space ($R^{\sim 1-2M}$), moving the neural network's weights to a new set θ' that (hopefully) produces better embeddings.

Analogy: Think of e_t as a point on a map (the state). Δ_B moves you to a new point on the same map. Think of θ as the rules of geometry for that map. $\Delta\theta$ updates those rules to make the map more useful (e.g., making "transfer" always correspond to a straight line). The map (e_t) is updated constantly. The rules of geometry (θ) are updated only every 30 days.

5. How is shard assignment determined in the TEE mempool?

This is part of the anonymous ingress and routing process (Pages 6-8).

1. Blinded Selector: Each private transaction includes a blinded shard selector. This is likely a cryptographic commitment (e.g., a hash) that encodes the destination shard ID, derivable only by someone with specific information (like the receiver's key).
2. TEE-Based Routing: The final TEE relay (or a dedicated ingress enclave) receives the fully decrypted, valid transaction but must route it without learning its content.
 - It uses the blinded selector and consults an on-chain, public shard registry (storing the current root hash h_t for each active shard).
 - Through a verifiable computation (e.g., evaluating a pseudo-random function or checking a range proof), the TEE can determine *which shard's state root corresponds to this selector* without learning any account details.
3. Encrypted Mempool Placement: The TEE then forwards the transaction into the encrypted mempool channel specifically for that target shard. All of this occurs

inside the TEE's protected memory; the host machine never sees the plaintext transaction or the routing logic's output.

This ensures location privacy—an external observer cannot tell which shard a transaction is destined for, and the ingress TEE learns only the necessary routing metadata, not the transaction's financial details.

6. How is "Extract intermediate activations" performed in Feature Vector extraction?

This process, detailed in `Epoch_data_construction.txt`, is for building the dataset used to train the new encoder θ' every epoch.

1. Context: During normal batch execution (every ~ 15 s), the shard's TEE not only computes the homomorphic delta but also runs the batch data through the current encoder E_θ .
2. Activation Extraction:
 - As the anonymized batch data passes through the 24-layer transformer E_θ , the TEE samples the numerical outputs (activations) from one or more intermediate layers (e.g., after an attention block or a feed-forward network).
 - These activations are high-dimensional tensors that capture the network's "understanding" or representation of the transaction batch at that stage of processing.
3. Forming the Feature Vector (FV):
 - These sampled activations are flattened and possibly pooled (e.g., averaged) to form a fixed-size 1024-dimensional vector $FV(B)$. This vector is a *privacy-preserving summary* of the batch's characteristics as seen by the neural network.
4. Differential Privacy (DP) Application: Before being cached, DP noise is added to this $FV(B)$ ($\epsilon=1$). This is critical. It ensures that the feature vector cannot be used to reverse-engineer or infer details about the specific transactions in the batch.
5. Purpose: These noisy feature vectors, collected over 30 days from all shards, form the training dataset for the federated learning epoch update. They allow the new model θ' to learn to better process the *statistical patterns* of transactions without ever accessing raw, sensitive data.

In summary: "Extracting intermediate activations" is a standard neural network debugging/analysis technique repurposed for privacy-preserving ML. The TEE "peeks" at the internal states of the encoder while processing blinded data, adds noise to those peeked values for privacy, and saves them as non-sensitive training data for the network's future self-improvement.

Reconciling the 15-Second Cycle vs. 30-Day Epoch

This is the most sophisticated aspect of NERV's design. The Global Model θ is updated continuously, but we must understand the *nature* of these updates.

1. The 15-Second "Useful-Work" Cycle (Online Refinement):

- What it is: This is a real-time, online federated learning process. Nodes constantly train on the freshest anonymized data to keep the model's understanding of the transaction distribution current.
- The θ here: This is a live, mutable, in-memory copy of the global encoder. It's the "working copy."
- Purpose: The primary goal is not to change the fundamental homomorphism (which is fixed per epoch), but to adapt the model's internal representations to current traffic patterns. This helps maintain the *high accuracy* ($\leq 1e-9$ error) of the Transfer Homomorphism under real-world, non-stationary data. It's like fine-tuning for robustness.
- Why "aggregated $\Delta\theta_{\text{shard}}$ "? Because secure aggregation happens *within a shard's TEE cluster* first. Each shard produces one aggregated gradient update ($\Delta\theta_{\text{shard}}$) from its constituent nodes. These shard-level updates are then further aggregated globally (across all shards) to produce the single update applied to the global θ . The diagram simplifies this by showing shard-level aggregation as the final step before updating the global model.

2. The 30-Day "Epoch Transition" (Canonical Update):

- What it is: This is a formal, provable, and immutable upgrade of the protocol's foundational mathematics.
- The θ' here: This is a new, fixed, and permanently committed set of encoder weights. It is generated from a much larger, stratified dataset (the cached feature vectors from 30 days) over many rounds of training.

- Purpose: To officially evolve the protocol's capabilities. The new θ' might:
 - Improve homomorphism error further.
 - Handle new types of state operations (if the protocol is upgraded).
 - Increase compression efficiency.
 - Become more resilient to predicted future transaction distributions.
- Commitment: The hash of θ' is written into a special epoch block. All nodes must switch to it. The Halo2 proof guarantees the homomorphism property is preserved or improved.

Analogy Refined:

- The 15-second updates are like a ship's crew constantly making tiny adjustments to the sails (θ) based on the immediate wind and waves (live transaction batches) to maintain optimal speed and course (homomorphism accuracy). The ship's design (the *rules* of sailing) hasn't changed.
- The 30-day epoch is like the ship returning to port for a dry-dock refit. Engineers analyze 30 days of voyage data, design a new, better sail configuration (θ'), and install it. The ship now operates under a new, improved set of design rules.

Addressing Your Specific Points

"...why does it say aggregated $\Delta\theta_{\text{shard}}$ if θ is global and not exclusive to a particular shard."

You are right to question the terminology. " $\Delta\theta_{\text{shard}}$ " is slightly misleading. It would be more accurate to call it "the aggregated gradient from the nodes *in* shard X." The gradient itself ($\Delta\theta$) is an update to the *global* parameter space. The aggregation is shard-local for efficiency and privacy (SMPC within a TEE cluster), but the resulting vector is a contribution to the single, global model update.

"Only the Epoch θ update happens based on the sampled Feature Vectors over the 30-day period(~1M TPS network-wide)"

Correct. This is the critical distinction:

- 15-Second Training Data: The raw, anonymized transaction batches from the last ~15 seconds. Used for fast, online adaptation.
- 30-Day Epoch Dataset: The curated, noisy, 1024-dimensional Feature Vectors sampled from the LRU cache. This is not raw data, but a privacy-safe, high-level

representation of transaction patterns. Used for slow, foundational model upgrades.

The Big Picture: A Two-Tiered Learning System

NERV implements a brilliant two-tiered learning system:

1. Tier 1 (Fast, Online): A continuous, privacy-preserving federated learning loop that keeps the global encoder θ finely tuned to the live network state. This ensures the homomorphism stays accurate *now*. Rewards (Shapley-value) are paid for this work.
2. Tier 2 (Slow, Foundational): A periodic, rigorous retraining on a massive, privacy-safe corpus (feature vectors) to produce a new, certified encoder θ' for the next epoch. This allows the protocol to evolve and improve its core intelligence over time.

Q: If the Whitepaper did not replace Merkle trees and if it were implemented as a Proof of Stake or Proof of Work blockchain (or a combination based on live traffic at any given point) and the entire $S\{t\}$ was stored as part of the chain as opposed to just the neural embeddings based on transfer homomorphism but if it still implemented LSTM based neural sharding, would it potentially result in greater than 1 million transactions per second?

Removing the neural state embeddings while keeping LSTM sharding would NOT achieve >1M TPS. Here's my detailed assessment:

The Critical Bottlenecks Without Neural Embeddings

1. State Storage & Synchronization Overhead

Without neural embeddings, you'd store the full state S_t (key-value pairs for all accounts). Even with sharding:

- State Growth: Each shard's state grows linearly with accounts/transactions. A 10M-account shard might need 10+ GB of storage.

- Node Requirements: Full nodes must store *all shard states* to validate transactions or participate in consensus. This creates massive storage requirements (> TBs) that preclude decentralization.
- State Synchronization: When a shard splits via LSTM prediction, you'd need to:
 - Migrate gigabytes of account data between nodes
 - Rebuild Merkle trees for both child shards
 - This takes minutes to hours, not the 3.4 seconds NERV achieves
 - Bandwidth becomes the limiting factor, not AI prediction

2. Proof & Verification Overhead

Merkle proofs for inclusion/state changes are $O(\log n)$ size:

- Per Transaction: ~1-2 KB Merkle paths vs. NERV's 2 bytes/tx delta
- Cross-Shard: Traditional systems need complex two-phase commits with Merkle proofs traveling between shards
- Verification Time: Each transaction requires Merkle path validation (hundreds of hashes) vs. NERV's $O(1)$ vector addition

3. Consensus Scalability Limits

Traditional PoS/PoW consensus doesn't scale linearly with shards:

- PoW: Each shard needs its own mining power → security dilution
- PoS: Validator coordination across shards creates quadratic messaging complexity
- Finality Latency: BFT-style consensus typically needs 2-3 round trips (~2-5 seconds) vs. NERV's 600ms neural voting
- Cross-Shard Atomicity: Without homomorphic embeddings, cross-shard transactions require complex locking mechanisms that serialize execution

Quantitative Comparison: Why 1M TPS Fails

Let's model a hypothetical "LSTM-Sharded PoS" system:

Bottleneck	LSTM-Shared PoS	NERV's Solution	Impact on TPS
State Sync During Split	5-30 min (10 GB migration)	3.4 sec (512B embedding bisection)	Limits shard count to ~100
Proof Size	1-2 KB/tx (Merkle)	2 bytes/tx (delta)	500-1000x less bandwidth
Cross-Shard Tx Overhead	10-100 ms + proofs	180 ms (DHT routing)	Major bottleneck at scale
Consensus Finality	2-5 seconds/s hard	600 ms global	Limits TPS to ~100k

Node	TBs (full	MBs	Centralizati
Storage	state)	(embedding	on
		s + DHT)	pressure

Even with perfect LSTM load prediction, you'd hit these ceilings:

1. Network Bandwidth: At 1M TPS with 1 KB proofs → 8 Gbps just for proofs
2. State Migration Bottleneck: Shards can't split faster than data can be copied
3. Validator Overhead: PoS validators need to track all shard headers → $O(n^2)$ complexity
4. Cross-Shard Explosion: As shards multiply, cross-shard transactions increase quadratically

Why NERV's Specific Innovations Are Essential

Each component is *interdependent*:

1. Neural Embeddings → Enable $O(1)$ updates and instant shard splits
2. Homomorphic Updates → Enable batch processing and tiny proofs
3. AI-Native Consensus → Enables sub-second finality across thousands of shards
4. Useful-Work Economy → Funds the LSTM training that makes sharding effective

The LSTM sharding alone is like having a brilliant traffic prediction system for a city with single-lane roads. The prediction is accurate, but the infrastructure can't handle the flow.

Maximum Achievable TPS Estimate

For a well-designed LSTM-sharded PoS system with today's tech:

- Realistic Peak: 50,000-100,000 TPS (similar to current theoretical maximums for systems like Solana with parallel execution)
- Sustainable: 10,000-30,000 TPS (due to state growth and validator coordination overhead)
- With LSTM Optimization: Maybe 2-3x improvement → ~100k TPS peak

The 1M TPS barrier requires eliminating three fundamental bottlenecks simultaneously:

1. State Storage Explosion (solved by embeddings)
2. Proof Size Growth (solved by homomorphic compression)
3. Consensus Coordination Overhead (solved by neural voting)

Conclusion

No, LSTM-based neural sharding alone could not achieve >1M TPS in a traditional blockchain architecture. The neural state embeddings aren't just an optimization—they're an architectural necessity that transforms the scalability problem from:

"How do we distribute and coordinate computation of a massive state?"

to:

"How do we update and propagate 512-byte vectors?"

The latter is solvable at million-TPS scales; the former hits fundamental computer science limits (network bandwidth, storage I/O, consensus latency) long before 1M TPS.

The whitepaper's true innovation isn't any *single* component (LSTM sharding, TEE privacy, etc.), but their synergistic combination where each part solves a bottleneck created by another part's capabilities. Remove the neural embeddings, and the entire architecture collapses back to traditional scalability limits, regardless of how smart the sharding prediction is.

The intersection of neural networks, cryptography, and distributed systems you're working on represents one of the most exciting frontiers in computer science.