

NERV V2.0: The Self-Evolving Blockchain

Via Neural Weight Oscillator Paradigm & Pure-Cryptographic Privacy

Version: 2.0

Date: August 8, 2026

Author: Rajeev Ragunathan

License: MIT/Apache 2.0

Abstract

NERV v2.0 is the first layer-1 blockchain to simultaneously deliver absolute privacy-by-default, infinite horizontal scalability, post-quantum security from genesis, and perpetual self-improvement—without the existential dependency on massive pre-trained AI models or ZK-ML circuit bloat.

Previous architectures (including NERV v1.01) relied on deep neural networks (e.g., 24-layer Transformers) to compress blockchain state into 512-byte Neural State Embeddings. While mathematically elegant, forcing a highly non-linear Transformer to approximate linear state updates (Transfer Homomorphism) required ~7.9 million ZK circuit constraints and massive federated pre-training. This represented a multi-year, multi-million-dollar ML engineering nightmare, contradicting the goal of immediate production readiness.

NERV v2.0 replaces this with the Neural Weight Oscillator (NWO) Paradigm. By utilizing a single-layer Perceptron optimized in real-time by the Adam algorithm and Huber Loss, NERV v2.0 achieves native, exact (0-error) Transfer Homomorphism out of the box. This reduces the ZK circuit from 7.9M constraints to ~50K, eliminates pre-training entirely, and allows the network to learn and self-correct on every single block.

Furthermore, V2.0 removes the dependency on Hardware Enclaves (TEEs) for privacy, replacing the trusted-hardware mixer with a Staked, Post-Quantum Sphinx Mixnet + Threshold Decryption Mempool. This achieves 100% privacy relying solely on mathematics, making the network immune to hardware side-channel attacks while preserving >1M+ TPS scalability.

1. Introduction & Motivation

The blockchain industry faces a quadrilemma: speed, privacy, quantum-security, and intelligence. Existing networks sacrifice one or more of these properties.

High-throughput chains (Solana, Sui) are fully transparent. Private chains (Monero, Zcash) are slow and lack quantum resistance. No existing system natively adapts its core protocol to shifting transaction distributions over time.

NERV v2.0 solves this by fusing three cutting-edge paradigms:

1. NWO-Driven State: Replacing static Merkle trees with a continuously learning, 512-byte linear latent vector.
2. Pure-Crypto Privacy: Eliminating trusted hardware in favor of lattice-based mixnets and threshold mempools.
3. Post-Quantum Foundations: Utilizing only NIST-standardized lattice cryptography (ML-KEM, Dilithium) from genesis.

1.1 The First-Principles Shift: The NWO Paradigm

The core bottleneck of marrying AI and Cryptography is non-linearity. Neural networks gain power through non-linear activation functions (GELU, Softmax). However, ZK proof systems (Halo2, Plonky2) require arithmetization over prime fields. Non-linearities require complex lookup tables and millions of constraints, making ZK-ML computationally intractable for real-time blockchain state compression.

The Neural Weight Oscillator (NWO)—originally developed for lightweight, real-time financial forecasting—demonstrates that genuine, adaptive machine learning does not require deep networks. A single-layer Perceptron, continuously optimized via Backpropagation, the Adam Optimizer, and Huber Loss, can dynamically adapt to complex time-series data in real-time.

The NWO Breakthrough for Blockchain:

A single-layer perceptron is mathematically linear: $f(x) = W * x + b$.

Because it is strictly linear, it is natively and perfectly homomorphic:

$$f(x + \Delta x) = W * (x + \Delta x) + b = (W * x + b) + W * \Delta x = f(x) + f(\Delta x)$$

Impact:

- Zero Pre-Training: We do not need to train a model to "approximate" homomorphism. A linear model possesses it by mathematical definition.

- Trivial ZK-ML: Proving a dot product ($\mathbf{w} * \mathbf{x}$) in a ZK circuit requires only basic multiplication and addition gates. Circuit complexity drops by 99%.
- Real-Time Evolution: The network learns continuously. Every block produces a target state. The perceptron measures the Huber Loss against the target and updates its weights via Adam. The network "gets smarter" every 400ms without 30-day Federated Learning overhead.

2. Neural State Embeddings (Linear / NWO-Style)

The canonical state of a NERV shard is represented not as a massive key-value Merkle Patricia Trie, but as a 512-byte embedding $\mathbf{e}_t \in \mathbb{R}^{64}$ (64 dimensions of 64-bit fixed-point numbers).

2.1 Architecture

- Inputs (Features): Blinded Account Commitments, Transaction Deltas, and Network Load metrics.
- Weights (θ): A $64 \times N$ weight matrix $\mathbf{W}$ and bias vector $\mathbf{b}$ of size 64. These weights are public, community-owned parameters that evolve over time.
- Encoding: $\mathbf{e}_t = \mathbf{W} * \mathbf{S}_t + \mathbf{b}$

2.2 The Transfer Homomorphism (Exact)

For a private transfer $\mathbf{tx}$, the delta vector $\delta(\mathbf{tx})$ is computed client-side by the wallet using the network's public weights:

$$\delta(\mathbf{tx}) = \mathbf{W} * \Delta S(\mathbf{tx}) + \mathbf{b}_{\mathbf{tx}}$$

The shard updates the canonical state via simple vector addition:

$$\mathbf{e}_{\{t+1\}} = \mathbf{e}_t + \delta(\mathbf{tx})$$

In V1.01, error bounds ($\leq 10^{-9}$) were required because the Transformer only *approximated* linearity. In V2.0, the error is exactly 0. The linear model guarantees perfect additive state updates.

Up to 10,000 private transactions can be batched into a single aggregated delta Δ_B (averaging 0.5 bytes/tx overhead due to the linear aggregation of deltas), enabling massive TPS vertical scaling before horizontal sharding is even required.

3. Pure-Cryptographic Privacy (No TEEs)

Previous designs relied on 5-Hop TEE mixers. TEEs (Intel SGX, AMD SEV) are vulnerable to side-channel attacks (Plundervolt, SGX-Step) and represent a "trusted hardware" assumption, which violates the cypherpunk mandate of 100% mathematical privacy.

V2.0 achieves 100% privacy, zero metadata leakage, and post-quantum resistance from first principles by replacing TEEs with a purely cryptographic onion-routing and threshold-encryption architecture.

3.1 PQ-Sphinx Packet Format

User wallets construct 5-hop onion packets using the Sphinx format (pioneered by Nym). However, all layered encapsulation uses ML-KEM-768 (Post-Quantum Key Encapsulation Mechanism). Sphinx provides provable packet indistinguishability—packets change size and headers at every hop, preventing timing and volume correlation attacks.

3.2 Staked Mixnet

Relays are financially staked. If a relay drops a packet or acts maliciously, cryptographic proofs result in slashing. Relays add exponential jitter ($\mu=100\text{ms}$, $\sigma=200\text{ms}$) and generate dummy cover traffic. (A localized NWO Perceptron is used by relays to mimic real traffic distributions, replacing the TEE-bound LSTM previously required).

3.3 Threshold Decrypted Mempool

The final relay does *not* decrypt the transaction. Instead, the inner payload is encrypted under a network-wide Distributed Key Generation (DKG) public key. Transactions sit in the mempool fully encrypted. No single node can read the mempool.

3.4 Decryption for Execution

Only when a block producer assembles a batch of 10,000 encrypted txs do validators perform a threshold decryption ceremony. The txs are decrypted only within the validator's secure memory space for execution*, applied to the state embedding, and then immediately zeroed out.

* - This statement is incorrect. Please

Result: 100% privacy relying solely on Lattice-based encryption and Sphinx geometry.
No hardware trust assumptions. Immune to side-channels.

4. LatentLedger Lite: The ZK Circuit

Because the neural encoder is now a linear Perceptron, the ZK circuit is drastically simplified. We no longer need Halo2 lookup tables for GELU or Softmax non-linearities.

Circuit Objective: Prove that the client correctly computed $\delta(tx) = W * \Delta S(tx) + b_{tx}$ using the network's public weights W , without revealing the private transaction inputs $\Delta S(tx)$.

4.1 Circuit Steps (in Halo2 / Plonky2)

1. Witness Assignment: Assign private transaction deltas (sender, receiver, amount) as witness variables to advice columns.
2. Dot Product Gate: Enforce $y - W * x = 0$ using standard multiplication and addition custom gates. (Requires ~512 custom gates for a 64-dim embedding).
3. Addition Gate: Add bias b_{tx} (~512 addition gates).
4. Homomorphic Update Gate: Add $\delta(tx)$ to the public previous root e_t to get e_{t+1} .

4.2 Impact

- Constraint Count: Drops from 7.9 Million to ~50,000.
- Proof Time: Drops from ~1s (requiring GPU) to <50ms on a standard mobile CPU.
- Proof Size: Remains ~400–750 bytes (constant).
- Verification: <10ms on light clients.

This is production-ready today using off-the-shelf Halo2 or Plonky2 libraries. No ZK-ML research breakthroughs required.

5. Self-Evolution: Continuous Adam Backprop

The most brilliant aspect of the NWO is its use of the Adam Optimizer ($\beta_1=0.9$, $\beta_2=0.999$) and Huber Loss to continuously learn from errors. We apply this directly to the blockchain state, replacing 30-day Federated Learning rounds with per-block adaptation.

5.1 The Useful-Work Economy (Real-Time)

Every block, the network measures how well its current weights $\bar{w}$ are representing the state dynamics.

1. Target (y): The actual, agreed-upon canonical embedding root $e_{\{t+1\}}$ derived from the threshold-decrypted batch.
2. Prediction ($\hat{y}$): The output of the Perceptron using the previous state and the batch delta: $\hat{y} = \bar{w} * S_t + \bar{w} * \Delta_B + b$.
3. Huber Loss (L_δ):
 - If $|y - \hat{y}| \leq \delta$: $L_\delta(y, \hat{y}) = 0.5 * (y - \hat{y})^2$
 - If $|y - \hat{y}| > \delta$: $L_\delta(y, \hat{y}) = \delta * |y - \hat{y}| - 0.5 * \delta^2$

Why Huber? Blockchain state anomalies (e.g., massive exchange inflows, MEV spikes) act as outliers. Huber loss (unlike Mean Squared Error) prevents the network weights from violently overreacting to these spikes, ensuring state stability.

5.2 Adam Optimizer Update

Validators compute the gradients of the loss with respect to the weights $\bar{w}$ and bias b . They update the first moment (m) and second moment (v) estimates, apply bias correction, and adjust the weights:

1. $m = \beta_1 * m + (1 - \beta_1) * \text{gradient}$
2. $v = \beta_2 * v + (1 - \beta_2) * \text{gradient}^2$
3. $m_hat = m / (1 - \beta_1^{\text{step}})$
4. $v_hat = v / (1 - \beta_2^{\text{step}})$
5. $\bar{w}_{new} = \bar{w}_{old} - \alpha * m_hat / (\sqrt{v_hat} + \epsilon)$

5.3 Consensus on Weights

Validators include the hash of their computed $\bar{w}_{new}$ in their block votes. If $\geq 67\%$ of validators agree on the new weight hash (derived from applying Adam to the canonical state), the network's model weights are officially updated on-chain.

Result: The network literally evolves every block. If the transaction distribution shifts (e.g., network goes from DeFi heavy to NFT heavy), the Perceptron's weights shift to optimize compression and predictive accuracy for the new reality.

6. AI-Native Consensus & Dynamic Sharding

6.1 Consensus

Validators use their local NWO Perceptron to predict the next embedding hash: $h_{pred} = \text{Hash}(e_t + W * \Delta_B + b)$. Because this is a dot product, prediction takes microseconds. If $\geq 67\%$ of active stake agrees, instant probabilistic finality is reached in $< 600\text{ms}$. Disputes are resolved via Monte-Carlo simulation in validator memory.

6.2 Sharding

A secondary, dedicated NWO Perceptron handles dynamic sharding.

- Features: Current TPS, Mempool size, Inter-shard message queue, Gas limit.
- Target: Actual block execution time.
- Adaptation: The Perceptron learns which features best predict shard overload. When the Perceptron predicts load $> 90\%$, it triggers a deterministic bisection of the 512-byte embedding, splitting the shard into two new shards with re-initialized local weights. The Adam optimizer ensures the sharding model gets better at predicting network load over time.

7. Cryptographic Primitives Suite (Post-Quantum from Genesis)

Because 100% privacy is non-negotiable, we cannot compromise on post-quantum primitives. Grover's and Shor's algorithms mandate a lattice-based foundation.

- Signatures: Dilithium-3 (NIST standardized). Used for all validator attestations and transaction non-repudiation.
- Encryption/KEM: ML-KEM-768 (NIST standardized). Used for the 5-hop Sphinx mixnet and client-to-validator communication.

- ZK Proofs: Halo2 + Nova folding. (While we simplified the circuit, Halo2 remains ideal for recursive proving over the BLS12-381 curve, allowing infinite proof compression for our VDWs).
 - Aggregation: BLS12-381 used strictly for validator threshold signatures (fast multiparty signature aggregation for consensus).
-

8. Verifiable Delay Witnesses (VDWs)

The VDW is the user's permanent, offline-verifiable receipt of inclusion. Because the LatentLedger Lite circuit is so small, VDW generation drops from <150ms to <15ms per batch.

A VDW contains:

1. `tx_hash` (32 B)
2. `delta_proof` (~400 B, proving the linear dot product was correct via Halo2)
3. `embedding_root` (32 B)
4. `dilithium_sig` (2.4 KB)

Note: Dilithium signatures are larger than ECDSA, pushing VDW size to ~3.5 KB. This is an acceptable trade-off for Post-Quantum security. They are still tiny enough to be pinned to Arweave/IPFS and verified offline on a mobile device in <30ms.

9. Tokenomics & Genesis Allocation

NERV features a fixed supply of 10,000,000,000 (10 Billion) NERV. There is no perpetual inflation. Tokens are distributed at genesis to align long-term network incentives.

- Useful-Work Pool (72%): Allocated over time to validators who submit valid Adam gradients that reduce the network's global Huber Loss. This replaces traditional PoW/PoS staking rewards with AI-utility rewards.
- Visionary (6%): Originator allocation, vested linearly over 6 months.
- Code & Research (10%): Core codebase contributors, vested over 1 year with a 3-month cliff.

- Early Donors (8%): Hardware and audit capital providers, vested over 1 year with a 6-month cliff.
- Treasury (4%): Community-governed pool for grants and protocol upgrades.

10. Roadmap to Production (The Pragmatic Path)

By adopting the NWO Paradigm, NERV V2.0 bypasses the "Research Phase" entirely. We can generate production-ready code today using existing, battle-tested libraries.

Component	NERV V2.0 (NWO		Available Today?
	Traditional AI Approach	Paradigm)	
State Model	24-Layer Transformer	Single-Layer Perceptron	✔ (Basic Linear Algebra)
ZK Circuit	ZK-ML (7.9M constraints)	LatentLedger Lite (50K constraints)	✔ (Halo2 / Plonky2)
Privacy Mixer	5-Hop TEE (SGX/SEV)	5-Hop PQ-Sphinx + DKG Mempool	✔ (Nym + ML-KEM)

Self-Evolution	30-Day FL (Shapley Values)	Per-Block Adam / Huber Backprop	✓ (Standard PyTorch/Numpy logic)
----------------	----------------------------	---------------------------------	----------------------------------

PQ Crypto	Dilithium / ML-KEM	Dilithium / ML-KEM	✓ (liboqs / PQClean)
-----------	--------------------	--------------------	----------------------

Implementation Steps (Target: 10-Week Testnet)

1. Weeks 1-3: Circuit & State Engine. Implement the $w * x + b$ Halo2 circuit. Build the atomic State Transition Function (STF) and RocksDB persistence.
2. Weeks 4-6: P2P & Consensus. Integrate `rust-libp2p` with custom ML-KEM Noise handshakes. Implement BFT voting and staking/slashing logic.
3. Weeks 7-8: Privacy Layer. Implement the PQ-Sphinx relay logic and the DKG threshold mempool ceremony.
4. Weeks 9-10: AI & Launch. Wire Adam/Huber continuous backprop. Deploy 5-node testnet with Grafana monitoring.

11. Conclusion

NERV V2.0 retains the radical vision of the original architecture—a private, scalable, post-quantum, self-evolving financial nervous system. However, by shifting from the hubris of "we must build a massive ZK-Transformer" to the pragmatism of the Neural Weight Oscillator, we achieve a profound simplification.

We prove that true AI is not about static, pre-trained size; it is about continuous, localized adaptation. By making the model linear, we make the math perfectly homomorphic. By making the privacy purely cryptographic, we make it unbreakable. The code can be written today. The future is learned, private, and fast.

Appendix A: Cryptographic Primitives Suite

The NERV V2.0 architecture is built upon a strict "Post-Quantum from Genesis" mandate. The protocol entirely abandons classical Elliptic Curve Cryptography (ECC) and RSA for core consensus, privacy, and fund security, replacing them with NIST-standardized lattice-based cryptography. This appendix details the exact primitives utilized, their application within the protocol, and the rationale behind their selection.

A.1 Post-Quantum Public-Key Cryptography

1. CRYSTALS-Dilithium-3 (Digital Signatures)

- Standardization: NIST FIPS 204 (Module-Lattice-Based Digital Signature Algorithm).
- Where it is used:
 - Transaction non-repudiation (user spending keys).
 - Validator attestation signatures in block proposals and votes.
 - Signing Verifiable Delay Witnesses (VDWs) to create permanent, offline-verifiable inclusion receipts.
- Why it was chosen: Dilithium-3 provides a strong security level against both classical and quantum computers (Shor's algorithm). Unlike hash-based signatures (e.g., SPHINCS+), Dilithium produces reasonably sized signatures (~3.3 KB) with exceptionally fast verification times (<0.5ms). While larger than ECDSA signatures, the size is highly manageable within NERV's batched 10,000-transaction blocks and is a necessary tradeoff for quantum resistance.

2. ML-KEM-768 (Key Encapsulation Mechanism)

- Standardization: NIST FIPS 203 (Module-Lattice-Based Key Encapsulation Mechanism, formerly Kyber).
- Where it is used:
 - PQ-Sphinx Mixnet: Encapsulating ephemeral shared secrets for each of the 5 onion-routing hops.
 - Private Note Encryption: Encrypting shielded transaction outputs (private notes) for the recipient's wallet.
 - DKG Share Transport: Securely transmitting Distributed Key Generation shares between validators.

- Why it was chosen: ML-KEM-768 offers extremely fast key generation, encapsulation, and decapsulation. Its ciphertext size (~1088 bytes) is deterministic and small enough to be nested efficiently within the Sphinx packet format without causing network bloat. It provides the foundational IND-CCA2 security required for the threshold-encrypted mempool.

A.2 Threshold & Consensus Cryptography

3. BLS12-381 Pairing Cryptography

- Standardization: IETF draft-irtf-cfrg-pairing-friendly-curves.
- Where it is used:
 - Distributed Key Generation (DKG): The collective public key that encrypts the mempool is a BLS12-381 G1 point.
 - Threshold Decryption: Validators compute partial decryptions (G1 points) which are combined via Lagrange interpolation.
 - Consensus Aggregation: Block votes are signed using BLS signatures, allowing 67% quorum attestations to be aggregated into a single 48-byte signature.
- Why it was chosen: *Note on Post-Quantum status:* BLS12-381 is not post-quantum. However, in NERV V2.0, it is used strictly for liveness, mempool encryption, and consensus efficiency—not for long-term fund security. If a quantum computer were to break BLS, an attacker could forge consensus votes or decrypt the mempool, but they could not forge Dilithium-3 signatures to steal user funds. The use of BLS12-381 is a pragmatic engineering choice to achieve <600ms finality via signature aggregation; the protocol includes a hook to hard-fork to a PQ-aggregate signature scheme (e.g., lattice-based threshold signatures) before quantum hardware matures.

A.3 Zero-Knowledge Proof System

4. Halo2 + Nova Folding (Recursive ZK Proofs)

- Standardization: Privacy Scaling Explorations (PSE) Halo2 specifications.
- Where it is used: The LatentLedger Lite circuit. It proves that the client correctly computed the homomorphic state $\delta(tx) = W * \Delta S(tx) + b_{tx}$ using the

network's public weights w , without revealing the private transaction inputs

$\Delta S(tx)$.

- Why it was chosen: Halo2 operates over the BLS12-381 curve, allowing it to natively interoperate with the DKG threshold signatures and validator identity layer. The shift to the linear NWO Perceptron reduces the circuit constraints from 7.9 Million (V1.01's non-linear Transformer) to ~50,000. This allows proof generation in <50ms on mobile CPUs. Nova folding is layered on top to allow infinite recursive proof compression, batching 10,000 private transactions into a single ~400-byte proof.

A.4 Symmetric Cryptography & Hashing

5. BLAKE3

- Standardization: BLAKE3 specification.
- Where it is used:
 - Computing 32-byte transaction hashes ($TxHash$).
 - Hashing the 512-byte embedding vectors into 32-byte $EmbeddingRoots$.
 - Deriving symmetric keys from ML-KEM shared secrets (XOF).
 - Keyed hashing for Sphinx packet integrity.
- Why it was chosen: BLAKE3 is significantly faster than SHA-2 and SHA-3, operating at gigabytes-per-second on modern CPUs. It is a Merkle-tree friendly hash function natively, making it ideal for committing to the neural state embeddings. It also features a built-in Key Derivation Function (KDF) and Pseudorandom Function (PRF), reducing the need for external cryptographic crates.

6. SHA3-256 (Keccak)

- Standardization: NIST FIPS 202.
- Where it is used: Used in specific domain-separated contexts where standard NIST hashing is mandated for interoperability, such as mapping BLS messages to G2 points (hash-to-curve) and specific Dilithium challenge generation.
- Why it was chosen: Required for strict compliance with NIST FIPS 203/204 standards and BLS signature domain separation.

7. ChaCha20-Poly1305 (AEAD)

- Standardization: RFC 8439.
- Where it is used:
 - Encrypting the layers of the PQ-Sphinx onion packets.
 - Encrypting the plaintext transaction payloads inside the DKG threshold ciphertexts.
 - Encrypting wallet memos attached to private notes.
- Why it was chosen: ChaCha20-Poly1305 is an Authenticated Encryption with Associated Data (AEAD) cipher that is highly resilient to timing side-channel attacks. It is hardware-accelerated on almost all modern ARM and x86 processors, ensuring that the 5-layer Sphinx decryption process does not introduce latency for mixnet relays.

A.5 Key Derivation & Memory Hardness

8. PBKDF2-HMAC-SHA512 & BLAKE3-KDF

- Standardization: RFC 8018 (PBKDF2) and BLAKE3 KDF.
- Where it is used:
 - PBKDF2: Deriving the 64-byte master seed from the 256-bit BIP-39 mnemonic entropy (2048 iterations).
 - BLAKE3-KDF: Deriving hierarchical detection keys, diversified receiver commitments, and symmetric memo keys from the master seed.
- Why it was chosen: PBKDF2 is retained for BIP-39 mnemonic compatibility, ensuring users can recover funds using standard hardware wallet backups. BLAKE3-KDF is used for rapid, domain-separated internal key derivation, ensuring that a compromised detection key cannot be used to derive a spending key.

Summary of the Pragmatic Security Posture:

NERV V2.0 does not rely on untested, experimental cryptography. By combining the speed of BLAKE3 and ChaCha20, the quantum resistance of Dilithium and ML-KEM, the recursive efficiency of Halo2, and the aggregation power of BLS12-381, the protocol achieves a 100% pure-cryptographic privacy model that is mathematically sound and available for production deployment today.

Appendix B: End-to-End Transaction Lifecycle & Mathematical Formalism

This appendix provides a comprehensive, step-by-step walkthrough of a complete private transaction lifecycle on the NERV V2.0 network. It traces the flow from user intent to final cryptographic confirmation, detailing the exact mathematical operations, cryptographic primitives, and network protocols invoked at each stage. Note: Please refer to 'NERV Whitepaper V2.0 - Addendum' for an expanded version of this section.

Phase 1: Client-Side Transaction Construction

Let a user, Alice, wish to transfer v nano-NERV to Bob.

1.1 Private Note Selection & State Delta (ΔS)

Alice's wallet locally stores unspent private notes. She selects a set of notes $I = \{n_1, n_2, \dots\}$ such that $\sum_{i \in I} v_i \geq v + \text{fee}$. She constructs outputs: Bob's note n_{bob} and her change note n_{change} .

The local scalar state delta is calculated as:

$$\Delta S = v_{\text{bob}} + v_{\text{change}} - \sum_{i \in I} v_i$$

1.2 Neural State Embedding & Homomorphic Delta (δ_{tx})

The canonical state of the shard is represented as a 512-byte embedding vector $e_{\text{sh}} \in \mathbb{R}^{64}$ (64 dimensions of 64-bit fixed-point numbers). The network's public NWO

Perceptron weights are a matrix $W \in \mathbb{R}^{64 \times N}$ and bias $b \in \mathbb{R}^{64}$.

Because the perceptron $f(x) = W \cdot x + b$ is strictly linear, it is perfectly homomorphic.

Alice computes the transaction's homomorphic delta:

$$\delta_{\text{tx}} = W \cdot \Delta S + b_{\text{tx}}$$

where b_{tx} is a transaction-specific bias derived from the nullifiers to prevent replay attacks.

1.3 Zero-Knowledge Witness Assignment (LatentLedger Lite)

Alice must prove she computed δ_{tx} correctly without revealing ΔS . She assigns her private variables to a Halo2 circuit:

1. Dot Product Gate: Proves $y = W \cdot \Delta S$ using ~ 512 multiplication and addition custom gates.
2. Addition Gate: Proves $\delta_{\text{tx}} = y + b_{\text{tx}}$.

Homomorphic Update Gate: Proves $e_{\text{sh},+1} = e_{\text{sh}} + \delta_{\text{tx}}$.

3. The circuit generates a proof π of constant size (~ 400 bytes) requiring only $\sim 50,000$ constraints.

1.4 Threshold Encryption

To hide the transaction payload (nullifiers, commitments, and encrypted notes) from the mempool, Alice encrypts it using the network's Distributed Key Generation (DKG) collective public key PK_{DKG} , yielding a ThresholdCiphertext C_{tx} .

Phase 2: PQ-Sphinx Onion Construction & Mixnet Routing

To prevent network-level correlation, Alice wraps C_{tx} in a 5-hop Post-Quantum Sphinx packet.

2.1 Onion Construction

Alice selects 5 staked relays $\{R_1, R_2, R_3, R_4, R_5\}$. For each relay R_i , she uses ML-KEM-768 to encapsulate an ephemeral shared secret ss_i .

She derives symmetric keys $K_i = \text{BLAKE3}(ss_i)$ and encrypts the routing instructions and payload layer-by-layer using ChaCha20-Poly1305:

$$\text{Layer}_i = \text{Enc}_{K_i}(R_{i+1}, \text{Layer}_{i+1})$$

The final packet contains the outermost KEM ciphertext ct_1 and Layer_1 .

2.2 Relay Logic, Jitter, & Cover Traffic

Alice broadcasts the packet to R_1 via libp2p's Kademlia DHT.

- Peeling: R_1 decapsulates ss_1 , decrypts Layer_1 , learns R_2 's address, and forwards Layer_2 .
- Jitter: To defeat timing analysis, R_1 holds the packet for a delay sampled from an exponential distribution $\Delta t \sim \text{Exp}(\mu=100\text{ms}, \sigma=200\text{ms})$.
- Cover Traffic: R_1 generates dummy packets using a local NWO Perceptron to mimic real traffic distributions, mixing them with real packets to defeat volume correlation.

Phase 3: Mempool Integration & Gossip

R_5 (the final exit node) passes the unencrypted C_{tx} (which remains encrypted under PK_{DKG}) to a validator node.

The transaction enters the Threshold Decrypted Mempool. Validators gossip the encrypted transactions via Gossipsub. No single node can read the transaction contents. The mempool batches up to 10,000 encrypted transactions $\{C_{tx}^1, \dots, C_{tx}^{10000}\}$ into a batch B .

Phase 4: Threshold Decryption Ceremony

A block producer proposes B. To execute the state transition, $\geq t$ validators must collaboratively decrypt the batch.

4.1 Partial Decryptions

For each C_{tx} , a validator V_i uses their DKG secret share s_i to compute a partial decryption $P_i = C_1^{s_i}$ in the BLS12-381 G1 group.

4.2 Lagrange Interpolation

The validators broadcast their partial decryptions. Once t partials are collected, the block producer combines them using Lagrange coefficients λ_i :

$$\text{Shared Secret Point} = \prod_{i=1}^t (P_i^{\lambda_i}) = C_1^{(\sum s_i \lambda_i)} = C_1^s$$

where s is the full DKG secret. The producer derives the symmetric key $K = \text{Hash}(C_1^s)$ and decrypts the 10,000 plaintext transactions.

Phase 5: State Transition & NWO Execution

5.1 Delta Aggregation

The producer aggregates the homomorphic deltas of the 10,000 transactions. Due to linearity, this is a simple vector addition:

$$\Delta_B = \sum_{j=1}^{10000} \delta_{txj}$$

5.2 Embedding Update

The canonical neural state embedding is updated via exact vector addition:

$$e_{+1} = e + \Delta_B$$

The producer computes the new embedding root $h_{+1} = \text{BLAKE3}(e_{+1})$.

Phase 6: Self-Evolution (Adam & Huber Loss)

The network measures how well its weights W represent the state dynamics and evolves.

6.1 Huber Loss Calculation

The target y is the actual agreed-upon canonical embedding e_{+1} . The prediction $\hat{y}$ is the output of the Perceptron: $\hat{y} = W \cdot S + W \cdot \Delta_B + b$.

The Huber Loss L_δ is computed to measure the error, providing robustness to MEV outliers:

$$L_{\delta}(y, \hat{y}) = \begin{cases} \frac{1}{2}(y - \hat{y})^2 & \text{if } |y - \hat{y}| \leq \delta \\ \delta(|y - \hat{y}| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

(where $\delta = 1.0$).

6.2 Adam Optimizer & Gradient Validation

Validators compute the gradient of the loss with respect to the weights: $g = \nabla_W L_{\delta}$.

Gradients are clipped to a max norm to prevent adversarial updates. The Adam optimizer updates the first moment m and second moment v :

$$m = \beta_1 m + (1 - \beta_1)g$$

$$v = \beta_2 v + (1 - \beta_2)g^2$$

Applying bias correction, the new weights are:

$$W_{\text{new}} = W - \alpha \times [m / (1 - \beta_1^t)] / [\sqrt{(v / (1 - \beta_2^t)) + \epsilon}]$$

Validators who submit gradients that successfully reduce L_{δ} are eligible for Useful-Work Rewards from the emission pool.

Phase 7: AI-Native Consensus & Finality

7.1 Perceptron Hash Prediction

Because the math is a simple dot product, validators predict the next embedding hash in microseconds:

$$h_{\text{pred}} = \text{Hash}(e_{\text{next}} + W \cdot \Delta_B + b)$$

7.2 Weighted Quorum

Validators vote on h_{pred} via BLS signature aggregation. The voting weight is $W_{\text{vote}} = \text{Stake} \times \text{Reputation}$.

Finality is achieved probabilistically in <600ms if:

$$\sum W_{\text{vote}} \geq 0.67 \times W_{\text{total}}$$

7.3 Monte Carlo Dispute Resolution

If a dispute arises (e.g., h_{pred} mismatch), validators run a Monte Carlo simulation in memory. They randomly select 32 validators and run 10,000 parallel simulations of the state transition to identify the malicious node. The loser of the dispute is slashed by 0.5% - 5% of their stake.

Phase 8: Dynamic Sharding & Erasure Coding

8.1 NWO Load Prediction

A secondary NWO Perceptron continuously evaluates network load using features $X = [\text{TPS}, \text{Mempool Size}, \text{Queue Depth}, \text{Gas Limit}]$. If the predicted load $L_{\text{pred}} = W_{\text{shard}} \cdot X > 0.90$, a shard split is triggered.

8.2 Embedding Bisection

The 64-dimension embedding e_{i+1} is deterministically bisected into two new 512-byte vectors e_A and e_B (e.g., interleaved or by halves), each inheriting re-initialized local weight matrices W_A and W_B .

8.3 Reed-Solomon Erasure Coding

Cross-shard messages and block data are encoded using Reed-Solomon with parameters $k=5$ (data shards) and $m=2$ (parity shards). This ensures that data remains available even if 2 out of 7 nodes in a shard go offline.

Phase 9: Receipt & User Confirmation

9.1 Verifiable Delay Witness (VDW) Generation

Once the block is finalized, a VDW is generated for Alice's transaction. The VDW (~3.5 KB) contains:

1. tx_hash (32 bytes)
2. π (Halo2 proof, ~400 bytes)
3. h_{i+1} (New embedding root, 32 bytes)
4. BLS threshold signature (48 bytes)
5. Dilithium-3 validator signature (3,293 bytes)

9.2 Offline Verification

The VDW is gossiped back to Alice's light wallet via libp2p Gossipsub. Alice verifies it offline in <30ms by:

1. Checking the BLS threshold signature against PK_DKG.
2. Checking the Dilithium-3 signature.
3. Running the Halo2 verifier on π .
4. Verifying homomorphic root reconciliation: $\text{Hash}(\text{trusted_prev_root} + \delta_{\text{tx}}) == h_{i+1}$.

9.3 Economic Settlement

The transaction is permanently confirmed. Alice's wallet updates its local balance, Bob's wallet trial-decrypts his new private note, and the validators receive their block rewards and Useful-Work gradient emission.

as currently written in Appendix B (Phase 4.2), the whitepaper has a massive privacy vulnerability because it allows the block producer to decrypt the individual transactions. Furthermore, the language regarding "secure memory space" and "immediate zeroing out" without a TEE is cryptographically unenforceable.

To make NERV V2.0 truly 100% private without hardware trust assumptions, we must leverage the linear homomorphic properties of the NWO Perceptron alongside threshold cryptography.

Here is a clean, structured breakdown addressing all your questions and providing the architectural fixes required to guarantee 100% privacy.

Appendix C: Corrections

1. How do we guarantee the validator's memory space is fully secure without a TEE?

The Reality: We cannot. Without a TEE (like Intel SGX), regular CPU memory is vulnerable to cold-boot attacks, memory dumps by a malicious hypervisor, or OS-level compromises.

The Fix: We must remove the language stating that transactions are "decrypted within the validator's secure memory space." Instead of relying on hardware security, we must rely on Cryptographic Memory Safety. This means transactions are never fully decrypted into any single validator's RAM. The block producer and validators operate purely on homomorphically encrypted data or aggregate values. If the plaintext never exists in a single node's memory, there is nothing to dump or steal.

2. How does the Sphinx mixnet on its own ensure 100% privacy without TEEs?

The Sphinx mixnet does not provide 100% privacy on its own; it provides Network-Level Anonymity.

Sphinx guarantees that an external observer (or even the relays themselves) cannot link the source IP to the destination IP, defeating timing and volume correlation attacks via exponential jitter and cover traffic. However, Sphinx only protects the transaction in transit. Once the packet reaches the exit node and enters the mempool, the transaction is still encrypted under the network's DKG public key. Execution privacy is only achieved by combining the Sphinx mixnet with the Threshold Mempool and the Zero-Knowledge proofs (Halo2). It is this triad—Sphinx + DKG Threshold Encryption + ZK Proofs—that eliminates hardware trust.

3. How are block producers selected, and isn't that too much power for one node?

The Vulnerability: As written in Appendix B (Phase 4.2), the block producer decrypts the batch. If a single entity is chosen to propose a block, they effectively become a centralized decryption oracle.

The Fix:

1. **Leader Election:** Block producers must be selected via a cryptographically verifiable random function (VRF) weighted by stake and reputation, rotating frequently.
2. **Stripping Producer Power:** The block producer's power must be reduced strictly to ordering and aggregating encrypted transactions. They do not get to decrypt them.

4. How do we ensure that tx's are "immediately zeroed out"?

The Reality: You cannot mathematically guarantee software zeroes out memory (e.g., via `memset` or Rust's drop traits) if the underlying OS or hypervisor is compromised.

The Fix: Replace "immediately zeroed out" with: "Because transactions remain encrypted under the network's DKG key and only aggregate state deltas are processed, no individual plaintext transaction is ever exposed to the block producer or validators, eliminating the need for memory zeroing."

5. Would the block producer view transaction details after decryption? (The Lagrange Interpolation Fix)

The Vulnerability: In the current whitepaper text (Phase 4.2), the block producer collects t partial decryptions, reconstructs the symmetric key K via Lagrange interpolation, and decrypts all 10,000 plaintext transactions. This completely breaks privacy. The producer can read every sender, receiver, and amount.

The Fix: We must use Lagrange interpolation of partial decryptions to ensure no single node can view individual transactions. Here is how we restructure the execution phase to achieve this:

The New Threshold Execution Flow (Replacing Appendix B, Phase 4 & 5)

Instead of decrypting individual transactions, we exploit the exact linear homomorphism of the NWO Perceptron ($f(x) = W \cdot x + b$).

1. Client-Side Submission: Alice submits her encrypted payload C_{tx} containing her private state ΔS . She also submits a public ZK proof π proving she computed $\delta_{tx} = W \cdot \Delta S + b_{tx}$ correctly. (Note: The nullifier for preventing double-spends is submitted separately as a public hash, so it does not leak transaction metadata).
2. Producer Aggregation (No Decryption): The block producer batches 10,000 encrypted payloads $\{C_{tx}^1, \dots, C_{tx}^{10000}\}$. Because the DKG encryption scheme (BLS12-381) is linearly homomorphic, the producer can mathematically aggregate the ciphertexts to create a single batch ciphertext $C_{batch} = \sum (C_{tx}^i)$.
3. Distributed Decryption of the Aggregate: A committee of $\geq t$ validators uses their DKG secret shares s_i to compute partial decryptions P_i on the aggregate ciphertext C_{batch} .
4. Lagrange Interpolation for Aggregate Only: The block producer combines these partial decryptions using Lagrange coefficients to reveal only the aggregate symmetric key K_{batch} .
5. State Execution: The producer uses K_{batch} to decrypt the batch, yielding only the aggregate homomorphic delta $\Delta_B = \sum (\delta_{tx}^i)$. The producer applies this to the state embedding: $e_{t+1} = e_t + \Delta_B$.
6. Consensus: Validators verify the ZK proofs of the individual transactions and verify that the aggregate delta Δ_B matches the decrypted result. They then vote on the new state root h_{t+1} via BLS signature aggregation.

Summary of Necessary Whitepaper Text Revisions

To make NERV V2.0 truly 100% private, here are the essential updates to Section 3.4 and Appendix B (Phase 4 & 5):

- Update to 3.4: "Only when a block producer assembles a batch of 10,000 encrypted txs do validators perform a threshold decryption ceremony on the homomorphically aggregated batch ciphertext. Via Lagrange interpolation, validators collaboratively decrypt only the aggregate state delta (Δ_B). Individual transaction plaintexts are never decrypted by any single node. The aggregate delta is applied to the state embedding, achieving 100% privacy relying solely on lattice-based encryption, Sphinx geometry, and exact linear homomorphism."
- Update to Appendix B (Phase 4.2): Remove the step where the producer "derives the symmetric key K and decrypts the 10,000 plaintext transactions." Replace with: "The producer aggregates the ciphertexts homomorphically. Validators provide partial decryptions of the aggregate batch. The producer interpolates these to recover only the aggregate batch key, decrypting the single aggregate delta vector Δ_B ."

By implementing this structural change, we completely close the privacy loophole. The block producer mathematically cannot view individual transactions, and the network achieves true 100% cryptographic privacy without any reliance on TEEs or trusted hardware